



# CS 6820 – Machine Learning

Lecture 4 (Info only – not given in class)

Instructor: Eric S. Gayles, PhD.

Jan 16, 2018

# ML Tools

- Matlab
- GNU Octave
- SciPy, Numpy, SciKits
  - Recommend installing Jupyter from [www.anaconda.com](http://www.anaconda.com) (evolved from iPython)
  - Good Read–Eval–Print Loop (REPL) environment
  - <https://github.com/jupyter/jupyter/wiki/A-gallery-of-interesting-Jupyter-Notebooks>
- Basic Linear Algebra Subprograms (BLAS)
- LAPACK — Linear Algebra PACKage
- Shogun-toolbox
- MLlib – Apache
- Deeplearn.js
- ConvnetJS
- MLPack
- TensorFlow

# NumPy and SciPy

- NumPy and SciPy are open-source add-on modules to Python that provide common mathematical and numerical routines in pre-compiled, fast functions. These are growing into highly mature packages that provide functionality that meets, or perhaps exceeds, that associated with common commercial software like MatLab.
- The NumPy (Numeric Python) package provides basic routines for manipulating large arrays and matrices of numeric data.
- The SciPy (Scientific Python) package extends the functionality of NumPy with a substantial collection of useful algorithms, like minimization, Fourier transformation, regression, and other applied mathematical techniques.
- <https://docs.scipy.org/doc/scipy/reference/tutorial/>

# Math Essentials

- Machine learning algorithms execute numerical solutions from Optimization Theory (Calculus), Statistics, Linear Algebra, and Discrete Math.
- Linear Algebra
  - Useful for compact representation of linear transformations on data
  - Dimensionality reduction techniques
  - Solving systems of equations

# Math Essentials

- To have maximum effectiveness developing Machine Learning solutions, you will need good mathematical intuitions about certain general machine learning principles and algorithms.

# Math Essentials

- Choose the best algorithm for the problem
- Set parameter and initial conditions
- Estimate runtime for solutions
- Drive validation strategies and troubleshooting
- Recognize/avoid over- or underfitting
- Set confidence boundaries on solutions
- Choose the best coding approach, framework, language, etc.

# Math Essentials

- $a \in A$       *set membership:  $a$  is member of set  $A$*
- $| B |$       *cardinality: number of items in set  $B$*
- $\| \mathbf{v} \|$       *norm: length of vector  $v$*
- $\Sigma$       *summation*
- $\int$       *integral*
- $\mathbb{R}$       *the set of real numbers*
- $\mathbb{R}^n$       *real number space of dimension  $n$* 
  - $n = 2$  : plane or 2-space
  - $n = 3$  : 3- (dimensional) space
  - $n > 3$  :  $n$ -space or *hyperspace*

# Math Essentials

- $\mathbf{x}, \mathbf{y}, \mathbf{z},$   
 $\mathbf{u}, \mathbf{v}$  *vector* (bold, lower case)
- $\mathbf{A}, \mathbf{B}, \mathbf{X}$  *matrix* (bold, upper case)
- $y = f(x)$  *function (map)*: assigns unique value in range of  $y$  to each value in domain of  $x$
- $dy / dx$  *derivative* of  $y$  with respect to single variable  $x$
- $y = f(\mathbf{x})$  *function* on multiple variables, i.e. a vector of variables; *function* in  $n$ -space
- $\partial y / \partial x_i$  *partial derivative* of  $y$  with respect to element  $i$  of vector  $\mathbf{x}$



# Math Essentials

A *probability space* is a random process or experiment with three components:

- $\Omega$ , the set of possible *outcomes*  $O$ 
  - ◆ number of possible outcomes =  $|\Omega| = N$
- $F$ , the set of possible *events*  $E$ 
  - ◆ an event comprises 0 to  $N$  outcomes
  - ◆ number of possible events =  $|F| = 2^N$
- $P$ , the *probability distribution*
  - ◆ function mapping each outcome and event to real number between 0 and 1 (the *probability* of  $O$  or  $E$ )
  - ◆ probability of an event is *sum* of probabilities of possible outcomes in event

# Math Essentials

Given:

- A discrete random variable  $X$ , with possible values  $x = x_1, x_2, \dots, x_n$
- Probabilities  $p( X = x_i )$  that  $X$  takes on the various values of  $x_i$
- A function  $y_i = f( x_i )$  defined on  $X$

The *expected value* of  $f$  is the probability-weighted “average” value of  $f( x_i )$ :

$$E( f ) = \sum_i p( x_i ) \cdot f( x_i )$$

# Math Essentials

- Addition of two matrices

- matrices must be same size
- add corresponding elements:

$$c_{ij} = a_{ij} + b_{ij}$$

- result is a matrix of same size

$$\mathbf{C} = \mathbf{A} + \mathbf{B} =$$

$$\begin{pmatrix} a_{11} + b_{11} & \cdots & a_{1n} + b_{1n} \\ \vdots & \ddots & \vdots \\ a_{m1} + b_{m1} & \cdots & a_{mn} + b_{mn} \end{pmatrix}$$

- Scalar multiplication of a matrix

- multiply each element by scalar:

$$b_{ij} = d \cdot a_{ij}$$

- result is a matrix of same size

$$\mathbf{B} = d \cdot \mathbf{A} =$$

$$\begin{pmatrix} d \cdot a_{11} & \cdots & d \cdot a_{1n} \\ \vdots & \ddots & \vdots \\ d \cdot a_{m1} & \cdots & d \cdot a_{mn} \end{pmatrix}$$

# Math Essentials

The probability a discrete variable  $A$  takes value  $a$  is:  $0 \leq P(A=a) \leq 1$

Probability

Probabilities of alternative outcomes add:  $P(A \in \{a, a'\}) = P(A=a) + P(A=a')$

Alternatives

The probabilities of all outcomes must sum to one:  $\sum_{\text{all possible } a} P(A=a) = 1$

Normalization

$P(A=a, B=b)$  is the joint probability that both  $A=a$  and  $B=b$  occur.

Joint Probability

Variables can be “summed out” of joint distributions:

Marginalization

$$P(A=a) = \sum_{\text{all possible } b} P(A=a, B=b)$$

$P(A=a|B=b)$  is the probability  $A=a$  occurs given the knowledge  $B=b$ .

Conditional  
Probability  
Product Rule

$$P(A=a, B=b) = P(A=a) P(B=b|A=a) = P(B=b) P(A=a|B=b)$$

Bayes rule can be derived from the above:

Bayes Rule

$$P(A=a|B=b, \mathcal{H}) = \frac{P(B=b|A=a, \mathcal{H}) P(A=a|\mathcal{H})}{P(B=b|\mathcal{H})} \propto P(A=a, B=b|\mathcal{H})$$


\* Thanks to A Yu

# Math Essentials

The derivative of  $f: R \rightarrow R$  is a function  $f': R \rightarrow R$  st.

$$f'(x) = \frac{df(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

if the limit exists.



# Math Essentials


- The gradient is an important concept from calculus in the context of machine learning.
- Gradients generalize derivatives to scalar functions of several variables.

$$\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix} \quad \text{i.e.} \quad [\nabla f]_i = \frac{\partial f}{\partial x_i}$$

- Definition: The gradient of  $f$ : in  $R^2 \rightarrow R$ :

It is a function  $\nabla f: R^2 \rightarrow R^2$  given by

$$\nabla f(x, y) := \left( \frac{\partial f}{\partial x} \quad \frac{\partial f}{\partial y} \right)^T$$



In the plane

- Definition: The gradient of  $f: R^n \rightarrow R$  is a function  $\nabla f: R^n \rightarrow R^n$  given by

$$\nabla f(x_1, \dots, x_n) := \left( \frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_n} \right)^T$$



# Math Essentials

The **Hessian** matrix of  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  is a matrix of second-order partial derivatives:



$$\nabla^2 f = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_d} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_d \partial x_1} & \cdots & \frac{\partial^2 f}{\partial x_d^2} \end{bmatrix} \quad \text{i.e.} \quad [\nabla^2 f]_{ij} = \frac{\partial^2 f}{\partial x_i \partial x_j}$$

# Basic Intro: Gradient Descent



Partial derivatives provide the slope to direct the solution toward minimization of the loss function.

- Pick initial solution
- Repeat until optimized {
  - Evaluate slope and minima conditionality
  - Pick dimension(s)
  - Move a small amount in the direction decreasing loss}

$$\begin{aligned}\theta_1 &:= \theta_1 - \alpha \frac{\partial J}{\partial \theta_1} \\ \theta_2 &:= \theta_2 - \alpha \frac{\partial J}{\partial \theta_2} \\ &\vdots \\ \theta_k &:= \theta_k - \alpha \frac{\partial J}{\partial \theta_k}\end{aligned}$$





# Basic Intro: Gradient Descent





\* A Beginners Tutorial for Machine Learning Beginners, Hao

# Basic Intro: Gradient Descent



# Basic Intro: Gradient Descent



saddle point


local max

# Basic Intro: Gradient Descent

- Simple concept: follow the gradient *downhill*
- Process:
  1. Pick a starting position:  $\mathbf{x}^0 = (x_1, x_2, \dots, x_d)$
  2. Determine the descent direction:  $-\nabla f(\mathbf{x}^t)$
  3. Choose a learning rate:  $\eta$
  4. Update your position:  $\mathbf{x}^{t+1} = \mathbf{x}^t - \eta \cdot \nabla f(\mathbf{x}^t)$
  5. Repeat from 2) until stopping criterion is satisfied
- Typical stopping criteria
  - $\nabla f(\mathbf{x}^{t+1}) \sim 0$
  - some validation metric is optimized

# Gradient descent optimization

- Problems:
  - Choosing step size
    - too small  $\rightarrow$  convergence is slow and inefficient
    - too large  $\rightarrow$  may not converge
  - Can get stuck on “flat” areas of function
  - Easily trapped in local minima



Picture credit: Andrew Ng, Stanford University, Coursera Machine Learning